

Activity_Course 2 TikTok project lab

January 27, 2024

1 TikTok Project

Course 2 - Get Started with Python

Welcome to the TikTok Project!

You have just started as a data professional at TikTok.

The team is still in the early stages of the project. You have received notice that TikTok's leadership team has approved the project proposal. To gain clear insights to prepare for a claims classification model, TikTok's provided data must be examined to begin the process of exploratory data analysis (EDA).

A notebook was structured and prepared to help you in this project. Please complete the following questions.

2 Course 2 End-of-course project: Inspect and analyze data

In this activity, you will examine data provided and prepare it for analysis.

The purpose of this project is to investigate and understand the data provided. This activity will:

1. Acquaint you with the data
2. Compile summary information about the data
3. Begin the process of EDA and reveal insights contained in the data
4. Prepare you for more in-depth EDA, hypothesis testing, and statistical analysis

The goal is to construct a dataframe in Python, perform a cursory inspection of the provided dataset, and inform TikTok data team members of your findings. *This activity has three parts:*

Part 1: Understand the situation * How can you best prepare to understand and organize the provided TikTok information?

Part 2: Understand the data

- Create a pandas dataframe for data learning and future exploratory data analysis (EDA) and statistical activities
- Compile summary information about the data to inform next steps

Part 3: Understand the variables

- Use insights from your examination of the summary data to guide deeper investigation into variables

To complete the activity, follow the instructions and answer the questions below. Then, you will use your responses to these questions and the questions included in the Course 2 PACE Strategy Document to create an executive summary.

Be sure to complete this activity before moving on to Course 3. You can assess your work by comparing the results to a completed exemplar after completing the end-of-course project.

3 Identify data types and compile summary information

Throughout these project notebooks, you'll see references to the problem-solving framework PACE. The following notebook components are labeled with the respective PACE stage: Plan, Analyze, Construct, and Execute.

4 PACE stages

- [Plan] (#scrollTo=psz51YkZVwtN&line=3&uniquifier=1)
- [Analyze] (#scrollTo=mA7Mz_SnI8km&line=4&uniquifier=1)
- [Construct] (#scrollTo=Lca9c8XON8lc&line=2&uniquifier=1)
- [Execute] (#scrollTo=401PgchTPr4E&line=2&uniquifier=1)

4.1 PACE: Plan

Consider the questions in your PACE Strategy Document and those below to craft your response:

4.1.1 Task 1. Understand the situation

- How can you best prepare to understand and organize the provided information?

Begin by exploring your dataset and consider reviewing the Data Dictionary.

Reading the data directory gives me a very good understanding of what the dataset should look like ideally. Also reading the description of the deliverables. And the emails sent to me regarding the project

4.2 PACE: Analyze

Consider the questions in your PACE Strategy Document to reflect on the Analyze stage.

4.2.1 Task 2a. Imports and data loading

Start by importing the packages that you will need to load and explore the dataset. Make sure to use the following import statements: * import pandas as pd

- import numpy as np

```
[3]: import pandas as pd
import numpy as np
```

```
[4]: # Load dataset into dataframe
data = pd.read_csv("tiktok_dataset.csv")
```

4.2.2 Task 2b. Understand the data - Inspect the data

View and inspect summary information about the dataframe by **coding the following**:

1. data.head(10)
2. data.info()
3. data.describe()

Consider the following questions:

Question 1: When reviewing the first few rows of the dataframe, what do you observe about the data? What does each row represent?

Question 2: When reviewing the data.info() output, what do you notice about the different variables? Are there any null values? Are all of the variables numeric? Does anything else stand out?

Question 3: When reviewing the data.describe() output, what do you notice about the distributions of each variable? Are there any questionable values? Does it seem that there are outlier values?

Then, load the dataset into a dataframe. Creating a dataframe will help you conduct data manipulation, exploratory data analysis (EDA), and statistical activities.

Note: As shown in this cell, the dataset has been automatically loaded in for you. You do not need to download the .csv file, or provide more code, in order to access the dataset and proceed with this lab. Please continue with this activity by completing the following instructions.

```
[36]: data.head(10)
```

```
[36]:   # claim_status  video_id  video_duration_sec  \
0    1          claim  7017666017              59
1    2          claim  4014381136              32
2    3          claim  9859838091              31
3    4          claim  1866847991              25
4    5          claim  7105231098              19
5    6          claim  8972200955              35
6    7          claim  4958886992              16
```

7	8	claim	2270982263	41
8	9	claim	5235769692	50
9	10	claim	4660861094	45

	video_transcription_text	verified_status	\
0	someone shared with me that drone deliveries a...	not verified	
1	someone shared with me that there are more mic...	not verified	
2	someone shared with me that american industria...	not verified	
3	someone shared with me that the metro of st. p...	not verified	
4	someone shared with me that the number of busi...	not verified	
5	someone shared with me that gross domestic pro...	not verified	
6	someone shared with me that elvis presley has ...	not verified	
7	someone shared with me that the best selling s...	not verified	
8	someone shared with me that about half of the ...	not verified	
9	someone shared with me that it would take a 50...	verified	

	author_ban_status	likes_per_view	comments_per_view	shares_per_view	\
0	under review	0.056584	0.000000	0.000702	
1	active	0.549096	0.004855	0.135111	
2	active	0.108282	0.000365	0.003168	
3	active	0.548459	0.001335	0.079569	
4	active	0.622910	0.002706	0.073175	
5	under review	0.521454	0.005516	0.185069	
6	active	0.647958	0.007258	0.258429	
7	active	0.001958	0.000020	0.000091	
8	active	0.409364	0.001088	0.042306	
9	active	0.183612	0.002727	0.072714	

	video_view_count	video_like_count	video_share_count	\
0	343296.0	19425.0	241.0	
1	140877.0	77355.0	19034.0	
2	902185.0	97690.0	2858.0	
3	437506.0	239954.0	34812.0	
4	56167.0	34987.0	4110.0	
5	336647.0	175546.0	62303.0	
6	750345.0	486192.0	193911.0	
7	547532.0	1072.0	50.0	
8	24819.0	10160.0	1050.0	
9	931587.0	171051.0	67739.0	

	video_download_count	video_comment_count
0	1.0	0.0
1	1161.0	684.0
2	833.0	329.0
3	1234.0	584.0
4	547.0	152.0
5	4293.0	1857.0

6	8616.0	5446.0
7	22.0	11.0
8	53.0	27.0
9	4104.0	2540.0

[37]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19382 entries, 0 to 19381
Data columns (total 15 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   #                                       19382 non-null  int64
1   claim_status                          19084 non-null  object
2   video_id                              19382 non-null  int64
3   video_duration_sec                   19382 non-null  int64
4   video_transcription_text             19084 non-null  object
5   verified_status                      19382 non-null  object
6   author_ban_status                   19382 non-null  object
7   likes_per_view                       19084 non-null  float64
8   comments_per_view                    19084 non-null  float64
9   shares_per_view                      19084 non-null  float64
10  video_view_count                     19084 non-null  float64
11  video_like_count                     19084 non-null  float64
12  video_share_count                    19084 non-null  float64
13  video_download_count                 19084 non-null  float64
14  video_comment_count                  19084 non-null  float64
dtypes: float64(8), int64(3), object(4)
memory usage: 2.2+ MB
```

[38]: data.describe()

```
[38]:
```

	#	video_id	video_duration_sec	likes_per_view \
count	19382.000000	1.938200e+04	19382.000000	19084.000000
mean	9691.500000	5.627454e+09	32.421732	0.276093
std	5595.245794	2.536440e+09	16.229967	0.173006
min	1.000000	1.234959e+09	5.000000	0.000000
25%	4846.250000	3.430417e+09	18.000000	0.130240
50%	9691.500000	5.618664e+09	32.000000	0.264037
75%	14536.750000	7.843960e+09	47.000000	0.398482
max	19382.000000	9.999873e+09	60.000000	0.666648

	comments_per_view	shares_per_view	video_view_count	video_like_count \
count	19084.000000	19084.000000	19084.000000	19084.000000
mean	0.000954	0.054860	254708.558688	84304.636030
std	0.001326	0.050597	322893.280814	133420.546814
min	0.000000	0.000000	20.000000	0.000000

25%	0.000098	0.014445	4942.500000	810.750000
50%	0.000455	0.039739	9954.500000	3403.500000
75%	0.001268	0.081864	504327.000000	125020.000000
max	0.010280	0.265956	999817.000000	657830.000000

	video_share_count	video_download_count	video_comment_count
count	19084.000000	19084.000000	19084.000000
mean	16735.248323	1049.429627	349.312146
std	32036.174350	2004.299894	799.638865
min	0.000000	0.000000	0.000000
25%	115.000000	7.000000	1.000000
50%	717.000000	46.000000	9.000000
75%	18222.000000	1156.250000	292.000000
max	256130.000000	14994.000000	9599.000000

`data.head(10)` The first few lines of data, makes light of few things. All the first five entries are claims. So that would be worth exploring whether there's some sort of bias, or there should be some sort of randomisation or other sorting and filtering. It's also worth noting that the first line has no `video_comments`, indicating that comments were disabled. There seems to be little correlation between how many times a video is viewed and shared.

`data.info()` The first apparent observation is the number of total rows and the non-null values. There seems to be a connection between 298 rows and null values. There is a mix of datatypes, both, ints, floats and objects. Objects for `“verified_status”` and `“author_ban_status”` could possibly be booleans instead since they'd take less space and memory? The number(`#`), could do with a more descriptive name, rather than a special character.

`data.describe()` `“video_like_count”`, `“video_share_count”`, `“video_download_count”`, `“video_comment_count”` seems like columns worth investigating to establish whether the 0 value is an outlier worth filtering out or they're all relevant. Also the range of values for these fields are very wide, and would also indicate there's something in the data obscuring the view. They also have means that are very close to the 75% percentile, further implying that the data in the current state is not giving the whole picture.

`“video_view_count”` has an average of 254708, but the less looking at the quantiles it suggests that a few videos are increasing the average. Is it worth using a median here to compare?

All the objects are missing due to not being possible to do numerical operations on them. But the 3 of them could be boolean instead, which would make it a lot easier to gain insight without compromising the data.

4.2.3 Task 2c. Understand the data - Investigate the variables

In this phase, you will begin to investigate the variables more closely to better understand them.

You know from the project proposal that the ultimate objective is to use machine learning to classify videos as either claims or opinions. A good first step towards understanding the data might therefore be examining the `claim_status` variable. Begin by determining how many videos there are for each different claim status.

```
[33]: print(data.groupby("claim_status")["claim_status"].count())
claim_num = (data.groupby("claim_status")["claim_status"].count()) / (data.
    ↳groupby("claim_status")["claim_status"].count().sum())

print(data.groupby("claim_status")["claim_status"].count().sum())

print((len(data)) == (data.groupby("claim_status")["claim_status"].count().
    ↳sum()))
print(claim_num * 100)
```

```
claim_status
claim      9608
opinion    9476
Name: claim_status, dtype: int64
19084
False
claim_status
claim      50.345839
opinion    49.654161
Name: claim_status, dtype: float64
```

There are rows missing their claim_status, as previously established. Apart from that, they're very equally split

Next, examine the engagement trends associated with each different claim status.

Start by using Boolean masking to filter the data according to claim status, then calculate the mean and median view counts for each claim status.

```
[35]: mask_claim = data['claim_status'] == "claim"
mask_opinion = data[('claim_status')] == "opinion"

data[mask_opinion]
data[mask_claim]
```

```
[35]:
```

	#	claim_status	video_id	video_duration_sec	\
0	1	claim	7017666017	59	
1	2	claim	4014381136	32	
2	3	claim	9859838091	31	
3	4	claim	1866847991	25	
4	5	claim	7105231098	19	
...	
9603	9604	claim	3883493316	49	
9604	9605	claim	4765029942	9	
9605	9606	claim	3513102998	27	
9606	9607	claim	9461481859	27	

9607 9608

claim 1622115206

16

	video_transcription_text	verified_status	\
0	someone shared with me that drone deliveries a...	not verified	
1	someone shared with me that there are more mic...	not verified	
2	someone shared with me that american industria...	not verified	
3	someone shared with me that the metro of st. p...	not verified	
4	someone shared with me that the number of busi...	not verified	
...	
9603	a colleague discovered on the radio a claim th...	not verified	
9604	a colleague discovered on the radio a claim th...	verified	
9605	a colleague discovered on the radio a claim th...	not verified	
9606	a colleague discovered on the radio a claim th...	not verified	
9607	a colleague discovered on the radio a claim th...	not verified	

	author_ban_status	likes_per_view	comments_per_view	shares_per_view	\
0	under review	0.056584	0.000000	0.000702	
1	active	0.549096	0.004855	0.135111	
2	active	0.108282	0.000365	0.003168	
3	active	0.548459	0.001335	0.079569	
4	active	0.622910	0.002706	0.073175	
...	
9603	active	0.625010	0.004574	0.073999	
9604	active	0.658297	0.004446	0.145060	
9605	under review	0.236556	0.000897	0.050011	
9606	active	0.278612	0.001393	0.058910	
9607	banned	0.195480	0.001627	0.058388	

	video_view_count	video_like_count	video_share_count	\
0	343296.0	19425.0	241.0	
1	140877.0	77355.0	19034.0	
2	902185.0	97690.0	2858.0	
3	437506.0	239954.0	34812.0	
4	56167.0	34987.0	4110.0	
...	
9603	737177.0	460743.0	54550.0	
9604	546987.0	360080.0	79346.0	
9605	885521.0	209475.0	44286.0	
9606	356747.0	99394.0	21016.0	
9607	114288.0	22341.0	6673.0	

	video_download_count	video_comment_count
0	1.0	0.0
1	1161.0	684.0
2	833.0	329.0
3	1234.0	584.0
4	547.0	152.0


```

...
9603      8119.0      3372.0
9604      4537.0      2432.0
9605      1210.0       794.0
9606      1163.0       497.0
9607       284.0       186.0

```

[9608 rows x 15 columns]

```
[25]: # What is the average view count of videos with "opinion" status?
```

```
(data[mask_opinion]["video_view_count"]).mean()
```

```
data[mask_opinion].describe()
```

```
#data[mask_claim].describe()
```

```
[25]:
```

	#	video_id	video_duration_sec	likes_per_view	\
count	9476.000000	9.476000e+03	9476.000000	0.0	
mean	14346.500000	5.622382e+09	32.359856	NaN	
std	2735.629909	2.530209e+09	16.281705	NaN	
min	9609.000000	1.234959e+09	5.000000	NaN	
25%	11977.750000	3.448802e+09	18.000000	NaN	
50%	14346.500000	5.611857e+09	32.000000	NaN	
75%	16715.250000	7.853243e+09	47.000000	NaN	
max	19084.000000	9.999835e+09	60.000000	NaN	

	video_view_count	video_like_count	video_share_count	\
count	9476.000000	9476.000000	9476.000000	
mean	4956.432250	1092.729844	217.145631	
std	2885.907219	964.099816	252.269583	
min	20.000000	0.000000	0.000000	
25%	2467.000000	289.000000	34.000000	
50%	4953.000000	823.000000	121.000000	
75%	7447.250000	1664.000000	314.000000	
max	9998.000000	4375.000000	1674.000000	

	video_download_count	video_comment_count
count	9476.000000	9476.000000
mean	13.677290	2.697446
std	16.200652	4.089288
min	0.000000	0.000000
25%	2.000000	0.000000
50%	7.000000	1.000000
75%	19.000000	3.000000
max	101.000000	32.000000

Question: What do you notice about the mean and media within each claim category? The

average views are much higher for opinions than claims 501029 vs 4956

Now, examine trends associated with the ban status of the author.

Use `groupby()` to calculate how many videos there are for each combination of categories of claim status and author ban status.

```
[34]: data.groupby(['claim_status', 'author_ban_status'])['video_id'].count()
```

```
[34]: claim_status  author_ban_status
      claim      active      6566
           banned      1439
           under review  1603
      opinion      active      8817
           banned       196
           under review   463
      Name: video_id, dtype: int64
```

Question: What do you notice about the number of claims videos with banned authors? Why might this relationship occur?

The claims category have much higher numbers in the category of banned and under review. They also have lower number of active users.

Continue investigating engagement levels, now focusing on `author_ban_status`.

Calculate the median video share count of each author ban status.

```
[ ]:
```

```
[41]: # What's the median video share count of each author ban status?
      data.groupby(['author_ban_status'])['video_share_count'].median()
```

```
[41]: author_ban_status
      active      437.0
      banned    14468.0
      under review  9444.0
      Name: video_share_count, dtype: float64
```

Question: What do you notice about the share count of banned authors, compared to that of active authors? Explore this in more depth.

The median `video_share_count` is much higher for banned users than active users.

Use `groupby()` to group the data by `author_ban_status`, then use `agg()` to get the count, mean, and median of each of the following columns: `* video_view_count * video_like_count * video_share_count`

Remember, the argument for the `agg()` function is a dictionary whose keys are columns. The values for each column are a list of the calculations you want to perform.

```
[40]: data.
      ↳groupby(['author_ban_status'])["video_view_count","video_like_count","video_share_count"].
      ↳agg(['count', 'mean', 'median'])
```

```
[40]:
```

	video_view_count			video_like_count	
	count	mean	median	count	
author_ban_status					
active	15383	215927.039524	8616.0	15383	
banned	1635	445845.439144	448201.0	1635	
under review	2066	392204.836399	365245.5	2066	

		video_share_count		
	mean	median	count	mean
author_ban_status				
active	71036.533836	2222.0	15383	14111.466164
banned	153017.236697	105573.0	1635	29998.942508
under review	128718.050339	71204.5	2066	25774.696999

	median
author_ban_status	
active	437.0
banned	14468.0
under review	9444.0

Question: What do you notice about the number of views, likes, and shares for banned authors compared to active authors? Banned users are more popular in view_count, like_count and video_share. Almost by double compared to active ones. Even the under review status, is more popular than the active ones. but by average and median.

Now, create three new columns to help better understand engagement rates: * likes_per_view: represents the number of likes divided by the number of views for each video * comments_per_view: represents the number of comments divided by the number of views for each video * shares_per_view: represents the number of shares divided by the number of views for each video

```
[6]: # Create a likes_per_view column
data.insert(7,"likes_per_view",(data['video_like_count'] /
↳data['video_view_count']))

# Create a comments_per_view column
data.insert(8,"comments_per_view",(data['video_comment_count'] /
↳data['video_view_count']))

# Create a shares_per_view column
data.insert(9,"shares_per_view",(data['video_share_count'] /
↳data['video_view_count']))
```

```
↳ -----
```

```
ValueError                                Traceback (most recent call↳  
↳last)
```

```
<ipython-input-6-5c8583dd377e> in <module>  
    1 # Create a likes_per_view column  
----> 2 data.insert(7,"likes_per_view",(data['video_like_count'] /↳  
↳data['video_view_count']))  
    3  
    4 # Create a comments_per_view column  
    5 data.insert(8,"comments_per_view",(data['video_comment_count'] /↳  
↳data['video_view_count']))
```

```
/opt/conda/lib/python3.7/site-packages/pandas/core/frame.py in↳  
↳insert(self, loc, column, value, allow_duplicates)  
    4412         if not allow_duplicates and column in self.columns:  
    4413             # Should this be a different kind of error??  
-> 4414             raise ValueError(f"cannot insert {column}, already↳  
↳exists")  
    4415         if not isinstance(loc, int):  
    4416             raise TypeError("loc must be int")
```

```
ValueError: cannot insert likes_per_view, already exists
```

Use `groupby()` to compile the information in each of the three newly created columns for each combination of categories of claim status and author ban status, then use `agg()` to calculate the count, the mean, and the median of each group.

```
[25]: data.  
↳groupby(['claim_status','author_ban_status'])['comments_per_view',"likes_per_view","shares_↳  
↳agg(['median','mean','count'])
```

```
[25]:
```

		comments_per_view		
		median	mean	count
claim_status	author_ban_status			
claim	active	0.000776	0.001393	6566
	banned	0.000746	0.001377	1439
	under review	0.000789	0.001367	1603
opinion	active	0.000252	0.000517	8817
	banned	0.000193	0.000434	196
	under review	0.000293	0.000536	463

		likes_per_view			shares_per_view \
		median	mean	count	median
claim_status	author_ban_status				
claim	active	0.326538	0.329542	6566	0.049279
	banned	0.358909	0.345071	1439	0.051606
	under review	0.320867	0.327997	1603	0.049967
opinion	active	0.218330	0.219744	8817	0.032405
	banned	0.198483	0.206868	196	0.030728
	under review	0.228051	0.226394	463	0.035027

		mean count	
claim_status	author_ban_status		
claim	active	0.065456	6566
	banned	0.067893	1439
	under review	0.065733	1603
opinion	active	0.043729	8817
	banned	0.040531	196
	under review	0.044472	463

Question:

How does the data for claim videos and opinion videos compare or differ? Consider views, comments, likes, and shares. Claim videos generate more response from the audience. Both in terms sharing, likes and comments.

4.3 PACE: Construct

Note: The Construct stage does not apply to this workflow. The PACE framework can be adapted to fit the specific requirements of any project.

4.4 PACE: Execute

Consider the questions in your PACE Strategy Document and those below to craft your response.

4.4.1 Given your efforts, what can you summarize for Rosie Mae Bradshaw and the TikTok data team?

Note for Learners: Your answer should address TikTok's request for a summary that covers the following points:

- What percentage of the data is comprised of claims and what percentage is comprised of opinions?
- What factors correlate with a video's claim status?
- What factors correlate with a video's engagement level?

The percentage of claim 50.34 % and opinion 49.65%

What what seems to be a correlation in the video's engagement is the notoriety. The Claims that come from users that have been banned or are under reviewd, seem to be causing a lot of user generated traffic, shares, likes and comments.

Congratulations! You've completed this lab. However, you may not notice a green check mark next to this item on Coursera's platform. Please continue your progress regardless of the check mark. Just click on the "save" icon at the top of this notebook to ensure your work has been logged.